

Unconstrained FPOP for Poisson Loss

GSoC 2026 — Medium Test

William Zhang

February 2026

github.com/williamzhang7792/gsoc2026-gfpop-william-zhang

Optimal Partitioning

Given count data y_1, \dots, y_n and penalty $\beta > 0$, find the segmentation minimizing

$$\sum_{k=1}^K \sum_{i \in S_k} (\mu_k - y_i \log \mu_k) + \beta (K - 1)$$

Each μ_k is set to the segment MLE \bar{y}_{S_k} .

(The constant $\log(y_i!)$ term is omitted—it does not depend on μ and cancels in optimization.)

Optimal Partitioning

Given count data y_1, \dots, y_n and penalty $\beta > 0$, find the segmentation minimizing

$$\sum_{k=1}^K \sum_{i \in S_k} (\mu_k - y_i \log \mu_k) + \beta (K - 1)$$

Each μ_k is set to the segment MLE \bar{y}_{S_k} .

(The constant $\log(y_i!)$ term is omitted—it does not depend on μ and cancels in optimization.)

Segment neighborhood (Segmentor3IsBack): best model for every $K = 1, \dots, K_{\max}$.

Optimal partitioning: best model for a single β , without computing all K .

From $O(n^2)$ to $O(n)$: The FPOP Idea

Standard DP: let $F_t = \min$ penalized cost over y_1, \dots, y_t .

$$F_t = \min_{0 \leq \tau < t} \{F_\tau + C(y_{\tau+1:t}) + \beta\} \quad \Rightarrow \quad O(n^2)$$

From $O(n^2)$ to $O(n)$: The FPOP Idea

Standard DP: let $F_t = \min$ penalized cost over y_1, \dots, y_t .

$$F_t = \min_{0 \leq \tau < t} \{ F_\tau + C(y_{\tau+1:t}) + \beta \} \implies O(n^2)$$

FPOP (Maidstone et al. 2016): keep cost as a *function* of the segment mean.

$$Q_t(x) = \min \left\{ \underbrace{Q_{t-1}(x)}_{\text{continue}}, \underbrace{\min_{x'} Q_{t-1}(x') + \beta}_{\text{switch}} \right\} + \ell(y_t, x)$$

where $x = \log \mu$ and $\ell(y, x) = e^x - yx$.

From $O(n^2)$ to $O(n)$: The FPOP Idea

Standard DP: let $F_t = \min$ penalized cost over y_1, \dots, y_t .

$$F_t = \min_{0 \leq \tau < t} \{ F_\tau + C(y_{\tau+1:t}) + \beta \} \implies O(n^2)$$

FPOP (Maidstone et al. 2016): keep cost as a *function* of the segment mean.

$$Q_t(x) = \min \left\{ \underbrace{Q_{t-1}(x)}_{\text{continue}}, \underbrace{\min_{x'} Q_{t-1}(x') + \beta}_{\text{switch}} \right\} + \ell(y_t, x)$$

where $x = \log \mu$ and $\ell(y, x) = e^x - yx$.

Notice that Q_t is **piecewise** — each piece inherits the $ae^x + bx + c$ structure from the Poisson loss. Candidates that fall entirely above the switch line get pruned \Rightarrow empirically $O(n)$.

Three Operations Per Step

1. **Flat line** at $K_{t-1} + \beta$, where $K_{t-1} = \min_x Q_{t-1}(x)$
`set_to_unconstrained_min_of`
2. **Min-envelope** of continue curve and flat switch line
`set_to_min_env_of`
3. **Add loss** $\ell(y_t, x) = e^x - y_t x$

Three Operations Per Step

1. **Flat line** at $K_{t-1} + \beta$, where $K_{t-1} = \min_x Q_{t-1}(x)$
`set_to_unconstrained_min_of`
2. **Min-envelope** of continue curve and flat switch line
`set_to_min_env_of`
3. **Add loss** $\ell(y_t, x) = e^x - y_t x$

Each Q_t is stored as a linked list of pieces.

Root-finding (Newton's method) determines where pieces cross.

Backtracking recovers the optimal segmentation from Q_n .

Constrained → Unconstrained

PeakSegOptimal: 2 states (up/down), $N \times 2$ cost array.

- ▶ `set_to_min_less_of`: enforces $\mu_k \leq \mu_{k+1}$
- ▶ `set_to_min_more_of`: enforces $\mu_k \geq \mu_{k+1}$

Constrained → Unconstrained

PeakSegOptimal: 2 states (up/down), $N \times 2$ cost array.

- ▶ `set_to_min_less_of`: enforces $\mu_k \leq \mu_{k+1}$
- ▶ `set_to_min_more_of`: enforces $\mu_k \geq \mu_{k+1}$

My change: **one** state, $N \times 1$ cost array.

- ▶ `set_to_unconstrained_min_of`: call `Minimize()`, emit a single constant piece

Constrained → Unconstrained

PeakSegOptimal: 2 states (up/down), $N \times 2$ cost array.

- ▶ `set_to_min_less_of`: enforces $\mu_k \leq \mu_{k+1}$
- ▶ `set_to_min_more_of`: enforces $\mu_k \geq \mu_{k+1}$


My change: **one** state, $N \times 1$ cost array.

- ▶ `set_to_unconstrained_min_of`: call `Minimize()`, emit a single constant piece

Everything else (min-envelope, root-finding, backtracking) stays unchanged.

Worked Example: The Data


$$\mathbf{y} = [3, 3, 10, 10], \quad \beta = 2.$$



Expected result: two segments $[3, 3 | 10, 10]$ with means $\hat{\mu}_1 = 3$, $\hat{\mu}_2 = 10$.

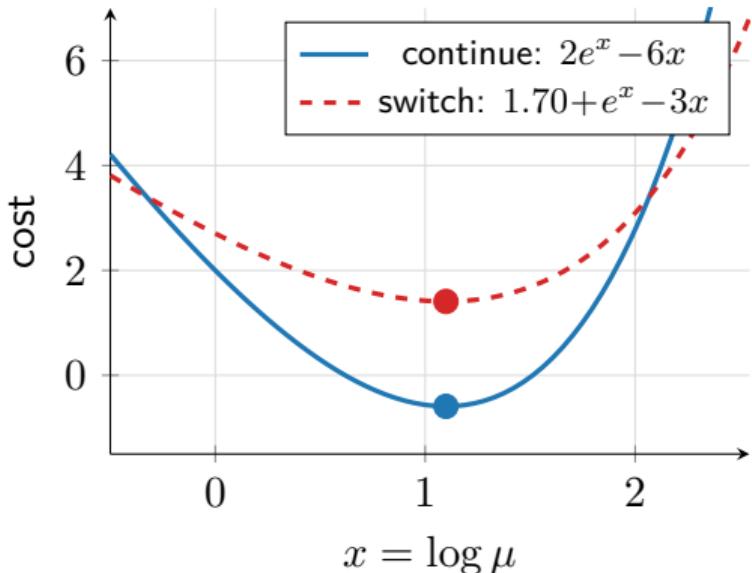
$t = 0$: Initial Cost ($y_0 = 3$)

$$Q_0(x) = e^x - 3x$$



Single piece. Minimum $K_0 \approx -0.30$ at $x = \log 3 \approx 1.10$.

$t = 1$: Continue vs. Switch ($y_1 = 3$)



Continue [3, 3]:

$\min \approx -0.59$ at $\mu = 3$

Switch (new seg [3]):

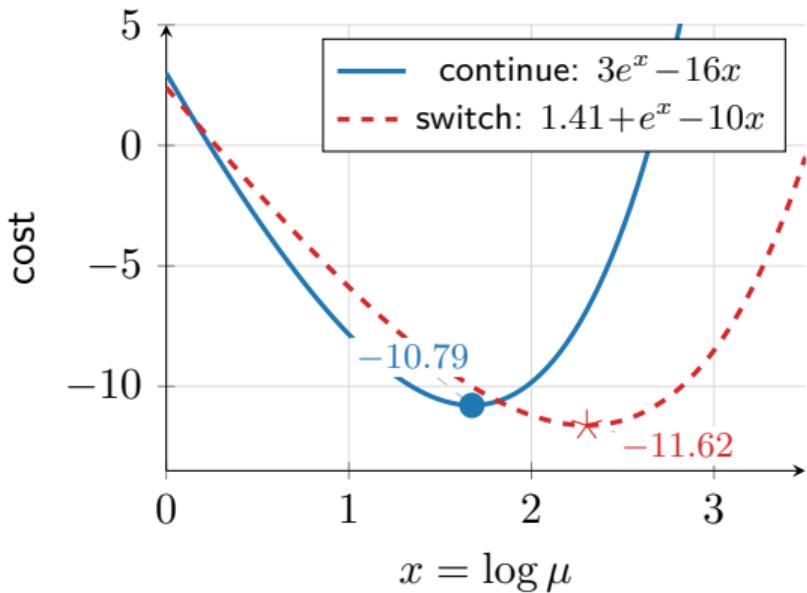
$\min \approx 1.41$ at $\mu = 3$

Continue wins by ≈ 2 .

Data is consistent—no reason to split.

$$K_1 \approx -0.59$$

$t = 2$: The Jump ($y_2 = 10$)



Continue [3, 3, 10]:

$\min \approx -10.79$ at $\mu \approx 5.3$

Switch [3, 3 | 10]:


$\min \approx -11.62$ at $\mu = 10$ *

Switch wins.

$y_2 = 10$ deviates far from $\bar{y}_{0:1} = 3$:
paying $\beta = 2$ to start fresh is worth it.

Curves cross at $x \approx 1.82$.

Result



Backtracking:

1. $t=3$: best mean $\mu=10$,
prev seg end = 1
2. $t=1$: best mean $\mu=3$,
prev seg end = -1 (start)

Two segments:

$$S_1 = \{y_0, y_1\}, \hat{\mu}_1 = 3$$

$$S_2 = \{y_2, y_3\}, \hat{\mu}_2 = 10$$

Penalized cost ≈ -24.64 .

Why Pruning is Exact

At step t , candidate τ is pruned if

$$Q_t(x, \tau) \geq K_t + \beta \quad \text{for all } x$$

i.e., continuing from τ is worse than starting fresh, for every possible mean.

Why Pruning is Exact

At step t , candidate τ is pruned if

$$Q_t(x, \tau) \geq K_t + \beta \quad \text{for all } x$$

i.e., continuing from τ is worse than starting fresh, for every possible mean.

Claim: a pruned candidate can never become optimal again.

Why Pruning is Exact

At step t , candidate τ is pruned if

$$Q_t(x, \tau) \geq K_t + \beta \quad \text{for all } x$$

i.e., continuing from τ is worse than starting fresh, for every possible mean.

Claim: a pruned candidate can never become optimal again.

1. At time t : candidate τ is replaced by $\tau' = t$, so $Q_t(x, \tau) > Q_t(x, \tau')$ for all x .
2. Future data adds the same loss to both:

$$Q_{t+s}(x, \tau) = Q_t(x, \tau) + L(x) > Q_t(x, \tau') + L(x) = Q_{t+s}(x, \tau')$$

3. Even if a future K_{t+s} is huge, τ still loses to τ' (which is alive).

Why Pruning is Exact

At step t , candidate τ is pruned if

$$Q_t(x, \tau) \geq K_t + \beta \quad \text{for all } x$$

i.e., continuing from τ is worse than starting fresh, for every possible mean.

Claim: a pruned candidate can never become optimal again.

1. At time t : candidate τ is replaced by $\tau' = t$, so $Q_t(x, \tau) > Q_t(x, \tau')$ for all x .
2. Future data adds the same loss to both:

$$Q_{t+s}(x, \tau) = Q_t(x, \tau) + L(x) > Q_t(x, \tau') + L(x) = Q_{t+s}(x, \tau')$$

3. Even if a future K_{t+s} is huge, τ still loses to τ' (which is alive).

⇒ Pruned = dead forever. FPOP is **exact**.

Pruning in Practice: Why $O(n)$

Consider a flat signal ($y_t = c$ for all t):

- ▶ The original candidate $\tau=0$ is the global optimum \rightarrow never pruned.
- ▶ Each new “switch” candidate $\tau=t$ pays penalty β but gains no advantage on flat data.
- ▶ As flat data accumulates, the false start’s cost rises above $K + \beta$.
- ▶ It gets pruned \rightarrow active candidate list stays small (often 2–3 pieces).

Pruning in Practice: Why $O(n)$

Consider a flat signal ($y_t = c$ for all t):

- ▶ The original candidate $\tau=0$ is the global optimum \rightarrow never pruned.
- ▶ Each new “switch” candidate $\tau=t$ pays penalty β but gains no advantage on flat data.
- ▶ As flat data accumulates, the false start’s cost rises above $K + \beta$.
- ▶ It gets pruned \rightarrow active candidate list stays small (often 2–3 pieces).

In general: the list grows when the signal changes, and shrinks (via pruning) when it stabilizes.

Amortized over the whole sequence: empirically $O(n)$.

Seeing a Pruning Step


At each step the active set holds one cost curve per surviving candidate.

The flat line at $K_t + \beta$ is the **ceiling**: if a candidate's curve sits entirely above it, that candidate is dead—it can never win for any future μ .

Candidate τ_1 dips below \rightarrow survives.

Candidate τ_2 is entirely above \rightarrow **pruned**.

On flat signals, new candidates are quickly pruned \rightarrow the active set stays small \rightarrow empirical $O(n)$.



Implementation

Starting from PeakSegOptimal:

- ▶ **Reused:** PoissonLossPieceLog, PiecewisePoissonLossLog, set_to_min_env_of, Minimize, root-finding, findMean
- ▶ **New:** set_to_unconstrained_min_of (7 lines—just Minimize + emit flat piece)
- ▶ **Simplified:** 2-state DP loop → 1-state, $N \times 2 \rightarrow N \times 1$
- ▶ **Added:** Rcpp wrapper, input validation, backtracking

Implementation

Starting from PeakSegOptimal:

- ▶ **Reused:** PoissonLossPieceLog, PiecewisePoissonLossLog, set_to_min_env_of, Minimize, root-finding, findMean
- ▶ **New:** set_to_unconstrained_min_of (7 lines—just Minimize + emit flat piece)
- ▶ **Simplified:** 2-state DP loop → 1-state, $N \times 2 \rightarrow N \times 1$
- ▶ **Added:** Rcpp wrapper, input validation, backtracking

Self-contained: PoissonFPOP.cpp (~760 lines, mostly reused infrastructure).

sourceCpp("PoissonFPOP.cpp") — no package install needed.

Validation

Compared against Segmentor3IsBack::Segmentor(model=1) using testthat.

8 test cases, 54 assertions:

| | |
|--------------------|-----------------------------------|
| 3-segment data | penalties 5, 10, 50, 100 |
| Single changepoint | penalties 10, 50, 200 |
| Constant-rate data | penalty 500 (1 segment) |
| 5-segment data | penalties 1, 5, 20, 100 |
| Data with zeros | penalties 1, 5, 50 |
| $\beta = 0$ | every point its own segment |
| $n = 2$ | 1-segment and 2-segment outcomes |
| Bad input | negative data, all-identical data |

All breakpoints and segment means match exactly.

Same results as Segmentor's $O(K_{\max} n^2)$ segment neighborhood, but in empirical $O(n)$.